

TASK – 1

Importing necessary libraries

import pandas as pd # Data manipulation

import numpy as np # Numerical computations

import matplotlib.pyplot as plt # Basic plotting

import seaborn as sns # Advanced visualization

print("Libraries imported successfully.")

import numpy as np

Creating Matrices

arr_zeros = np.zeros((3, 3))

print("Matrix of Zeros:\n", arr_zeros)

arr_ones = np.ones((2, 2))

print("Matrix of Ones:\n", arr_ones)

arr_full = np.full((1, 6), 5)

print("Full Matrix:\n", arr_full)

arr_eye = np.eye(4)

print("Identity Matrix:\n", arr_eye)

arr_range = np.arange(0, 14, 3)

print("Matrix Range:\n", arr_range)

arr_l = np.linspace(0, 6, 4)

print("Evenly distributed arrays:\n", arr_l)

print()

Arithmetic operations

arr = np.array([[1, 2, 3], [4, 5, 6]])

print("Addition:\n", arr + 2)

print("Multiplication:\n", arr * 5)

Universal functions

```
print("Square root:\n", np.sqrt(arr))
```

Shape and Reshape

```
print("Shape:", arr.shape)
```

```
print("Reshaped:\n", arr.reshape(1, 6))
```

Indexing and Slicing

```
print("Specific element:", arr[1, 2])
```

```
print("Conditional Selection:\n", arr[arr > 3])
```

Matrix of Zeros:

```
[[0. 0. 0.]
```

```
[0. 0. 0.]
```

```
[0. 0. 0.]]
```

Matrix of Ones:

```
[[1. 1.]
```

```
[1. 1.]]
```

Full Matrix:

```
[[5 5 5 5 5]]
```

Identity Matrix:

```
[[1. 0. 0. 0.]
```

```
[0. 1. 0. 0.]
```

```
[0. 0. 1. 0.]
```

```
[0. 0. 0. 1.]]
```

Matrix Range:

```
[ 0 3 6 9 12]
```

Evenly distributed arrays:

```
[0. 2. 4. 6.]
```

Addition:

```
[[3 4 5]
```

```
[6 7 8]]
```

Multiplication:

```
[[ 5 10 15]
```

```
[20 25 30]]
```

Square root:

```
[[1.      1.41421356 1.73205081]
```

```
[2.      2.23606798 2.44948974]]
```

Shape: (2, 3)

Reshaped:

```
[[1 2 3 4 5 6]]
```

Specific element: 6

Conditional Selection:

```
[4 5 6]
```

```
import pandas as pd
```

```
# Creating Series
```

```
data = [1, 2, 3, 4, 5]
```

```
serie = pd.Series(data)
```

```
# Custom index
```

```
index = ['a', 'b', 'c', 'd', 'e']
```

```
serie = pd.Series(data, index=index)
```

```
# DataFrame Creation
```

```
data = {
```

```
    "Name": ['Kitty', 'Bob', 'Alan'],
```

```
    "Age": [21, 20, 23],
```

```
    "City": ['LA', 'NY', 'CH']
```

```
}
```

```
df = pd.DataFrame(data)
print(df.head(1))
print(df.describe())
print(df.info())
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Simple Plot
```

```
x = [1, 2, 3, 4, 5]
y = [10, 20, 30, 40, 50]
plt.plot(x, y, marker='o')
plt.title("Sample Plot")
plt.show()
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Plotting the box plot
```

```
sns.boxplot(x="day", y="total_bill", data=tips)
plt.title("Box Plot of Total Bill by Day")
plt.show()
```

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
```

```
# Loading the dataset

df = pd.read_csv('Iris.csv')

print(df.head())


# Label Encoding

encoder = LabelEncoder()

df['Species'] = encoder.fit_transform(df['Species'])


# Standard Scaling

scaler = StandardScaler()

df_scaled = pd.DataFrame(scaler.fit_transform(df.drop(columns=['Id', 'Species'])))


# Histogram

plt.hist(df['SepalLengthCm'], bins=20, edgecolor='black')

plt.title('Sepal Length Distribution')

plt.xlabel('Sepal Length')

plt.ylabel('Frequency')

plt.show()


# Scatter Plot

sns.scatterplot(x='SepalLengthCm', y='SepalWidthCm', hue='Species', data=df)

plt.show()


# Heatmap for Correlation

corr_matrix = df.drop(columns=['Id']).corr()

sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')

plt.title('Correlation Heatmap')

plt.show()
```

TASK – 2

Import necessary libraries

import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler

from sklearn.decomposition import PCA

Step 1: Load the dataset

df = pd.read_csv('Iris.csv')

Step 2: Display first few rows and dataset info

print("DataFrame Head:\n", df.head())

print("DataFrame Info:\n", df.info())

Step 3: One-Hot Encoding for categorical columns

df_encoded = pd.get_dummies(df, columns=['Species'], drop_first=True)

print("Encoded DataFrame:\n", df_encoded.head())

Step 4: Standardize the features

scaler = StandardScaler()

x = df_encoded.drop(columns=['Id'])

x_scaled = scaler.fit_transform(x)

Step 5: Apply PCA for dimensionality reduction

pca = PCA(n_components=2)

x_pca = pca.fit_transform(x_scaled)

Step 6: Create PCA DataFrame and add 'Species' column

df_pca = pd.DataFrame(x_pca, columns=['PC1', 'PC2'])

df_pca['Species'] = df['Species']

```
# Step 7: Visualize original data
```

```
sns.scatterplot(x='SepalLengthCm', y='SepalWidthCm', hue='Species', data=df)
```

```
plt.title('Original Data: Sepal Length vs Sepal Width')
```

```
plt.xlabel('Sepal Length')
```

```
plt.ylabel('Sepal Width')
```

```
plt.show()
```

```
# Step 8: Visualize PCA results
```

```
sns.scatterplot(x='PC1', y='PC2', hue='Species', data=df_pca)
```

```
plt.title('PCA: Principal Component 1 vs Principal Component 2')
```

```
plt.xlabel('PC1')
```

```
plt.ylabel('PC2')
```

```
plt.show()
```

RFE

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.feature_selection import RFE
```

```
from sklearn.metrics import accuracy_score
```

```
# Load dataset
```

```
df = pd.read_csv('Iris.csv')
```

```
print(df.head())
```

```
print(df.info())
```

```
# Preprocessing
```

```
X = df.drop(columns=['Species'])
```

```
y = df['Species']

# Standardize and split data
X_scaled = StandardScaler().fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=40)

# Logistic Regression model
model = LogisticRegression(max_iter=10000)

# Evaluate accuracy before RFE
model.fit(X_train, y_train)
xpred=model.predict(X_test)
acc_b = accuracy_score(y_test, xpred)
print("Accuracy before RFE:", acc_b)

# Feature selection with RFE
rfe = RFE(model, n_features_to_select=2)
X_train_rfe = rfe.fit_transform(X_train, y_train)
X_test_rfe = rfe.transform(X_test)

# Print number of selected features
print("Number of selected features:", sum(rfe.support_))

# Evaluate accuracy after RFE
model.fit(X_train_rfe, y_train)
xpred=model.predict(X_test_rfe)
acc_a = accuracy_score(y_test, xpred)
print("Accuracy after RFE:", acc_a)

# Print selected features
print("Selected features:", X.columns[rfe.support_])
```



```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder, StandardScaler

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score


# Load and preprocess data

df = pd.read_csv('Pune_rent.csv')

df['price'] = pd.to_numeric(df['price'].str.replace(", ", ""))

df.fillna(df.mode().iloc[0], inplace=True)


# Encoding categorical features

label_encoder = LabelEncoder()

categorical_cols = df.select_dtypes(include=['object']).columns

for col in categorical_cols:

    df[col] = label_encoder.fit_transform(df[col])


# Feature selection

X = df[['bathroom', 'bedroom', 'area', 'locality', 'furnish_type']]

y = df['price']


# Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Scaling features

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)


# Train model

model = LinearRegression()
```

```
model.fit(X_train_scaled, y_train)
```

```
# Predictions and evaluation
```

```
y_pred = model.predict(X_test_scaled)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
mae = mean_absolute_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
# Display metrics
```

```
print(f'Mean Squared Error:{mse:.2f}')
```

```
print(f'Mean Absolute Error:{mae:.2f}')
```

```
print(f'Root Mean Squared Error:{np.sqrt(mse):.2f}')
```

```
print(f'R-Squared:{r2:.2f}')
```

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder, StandardScaler

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score


# Step 1: Load the dataset

df = pd.read_csv('User_Data.csv')


# Step 2: Data Preprocessing

# Convert 'Gender' column to numerical using Label Encoding

label_encoder = LabelEncoder()

df['Gender'] = label_encoder.fit_transform(df['Gender'])


# Fill missing values if any (Optional based on dataset)

df.fillna(df.mode().iloc[0], inplace=True)


# Step 3: Feature selection and target variable

X = df[['Gender', 'EstimatedSalary', 'Purchased']] # Independent variables

y = df['Age'] # Target variable


# Step 4: Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 5: Feature Scaling (Normalize the features)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)


# Step 6: Train the Linear Regression model

model = LinearRegression()
```

```
model.fit(X_train_scaled, y_train)
```

```
# Step 7: Make Predictions
```

```
y_pred = model.predict(X_test_scaled)
```

```
# Step 8: Evaluate Model Performance
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
mae = mean_absolute_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
# Step 9: Print Evaluation Metrics
```

```
print(f'Mean Squared Error (MSE): {mse:.2f}')
```

```
print(f'Mean Absolute Error (MAE): {mae:.2f}')
```

```
print(f'Root Mean Squared Error (RMSE): {np.sqrt(mse):.2f}')
```

```
print(f'R-Squared (R2): {r2:.2f}')
```

```
import pandas as pd

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix


# Step 1: Load the Iris dataset

iris = load_iris()

df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

df['species'] = iris.target # Add target variable


# Step 2: Prepare the features and target

X = df.drop('species', axis=1) # Features (input variables)

y = df['species'] # Target variable


# Step 3: Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 4: Create and train the model

model = DecisionTreeClassifier()

model.fit(X_train, y_train)


# Step 5: Make predictions

y_pred = model.predict(X_test)


# Step 6: Evaluate the model

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")

print("Classification Report:\n", classification_report(y_test, y_pred))

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
import pandas as pd

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix


# Step 1: Load the Iris dataset

iris = load_iris()

df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

df['species'] = iris.target # Add target variable


# Step 2: Prepare the features and target

X = df.drop('species', axis=1) # Features (input variables)

y = df['species'] # Target variable


# Step 3: Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Step 4: Create and train the KNN model

knn = KNeighborsClassifier(n_neighbors=3) # You can adjust the number of neighbors

knn.fit(X_train, y_train)


# Step 5: Make predictions

y_pred = knn.predict(X_test)


# Step 6: Evaluate the model

accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")

print("Classification Report:\n", classification_report(y_test, y_pred))

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Step 1: Load the dataset

df = pd.read_csv('loan_data_set.csv')


# Step 2: Prepare features and target

X = df.drop(columns=['Loan_Status'])

y = df['Loan_Status']


# Step 3: Encode categorical features

label_encoder = LabelEncoder()

X_encoded = X.apply(label_encoder.fit_transform)


# Step 4: Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)


# Step 5: Create and train the Random Forest model

rf = RandomForestClassifier()

rf.fit(X_train, y_train)


# Step 6: Make predictions

y_pred = rf.predict(X_test)


# Step 7: Evaluate the model

accuracy = accuracy_score(y_test, y_pred)

print(f"Random Forest Accuracy: {accuracy:.2f}")

print("Classification Report:\n", classification_report(y_test, y_pred))

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
from sklearn.datasets import load_iris

from sklearn.preprocessing import StandardScaler

from sklearn.cluster import KMeans

from sklearn.metrics import adjusted_rand_score, silhouette_score

import pandas as pd


# Load the dataset

iris = load_iris()

df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

df['species'] = iris.target


# Preprocessing: Standardize the data

scaler = StandardScaler()

x_scaled = scaler.fit_transform(df.drop(columns='species'))


# Apply K-Means Clustering

kmeans = KMeans(n_clusters=3, random_state=42)

kmeans.fit(x_scaled)

df['cluster'] = kmeans.labels_


# Performance Measures: Adjusted Rand Index and Silhouette Score

ari = adjusted_rand_score(df['species'], df['cluster'])

silhouette = silhouette_score(x_scaled, df['cluster'])


print("Cluster Centers:\n", kmeans.cluster_centers_)

print("Adjusted Rand Index:", ari)

print("Silhouette Score:", silhouette)
```



```
from scipy.cluster.hierarchy import dendrogram, linkage, fcluster
from sklearn.metrics import adjusted_rand_score, silhouette_score
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import load_iris
```

```
# Load the dataset
```

```
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
df['species'] = iris.target
```

```
# Preprocessing: Standardize the data
```

```
scaler = StandardScaler()
x_scaled = scaler.fit_transform(df.drop(columns='species'))
```

```
# Apply Hierarchical Clustering (Ward's method)
```

```
linkage_matrix = linkage(x_scaled, method='ward')
```

```
# Plot Dendrogram
```

```
dendrogram(linkage_matrix, labels=iris.target)
plt.title('Dendrogram')
plt.xlabel('Samples')
plt.ylabel('Distance')
plt.show()
```

```
# Get clusters
```

```
clusters = fcluster(linkage_matrix, t=3, criterion='maxclust')
df['cluster'] = clusters
```

```
# Performance Measures: Adjusted Rand Index and Silhouette Score
```

```
ari = adjusted_rand_score(df['species'], df['cluster'])  
silhouette = silhouette_score(x_scaled, df['cluster'])  
  
print("Cluster Distribution:\n", df.groupby(['cluster', 'species']).size())  
print("Adjusted Rand Index:", ari)  
print("Silhouette Score:", silhouette)
```